

OPERATIONS VIEW OF THE JAVA VM

Bill Schwanitz

Presentation available at
https://github.com/bilsch/columbus_devops_presentations

Rough overview


- Quick overview of java
- Crash course introduction to the Java heap
- Crash course on garbage collection (I'll keep this very simple)
- Common simple tunable/run-time tweaks
- Helpful tools
- Common exceptions
 - OutOfMemory
 - How to read/decrypt an exception
- “Just give me more heap!”
 - Developers like to punt and just ask for more heap. I'll give you compelling arguments to not give in too quickly

Quick overview of Java

- Object-oriented programming language developed by Sun Microsystems
 - First released in 1995
- Inspired/based on c and c++
- Original idea was to have a cross-platform just-in-time compiled language.
 - Write to the Java spec, it should run anywhere
 - The nasty stuff supposedly gets handled by the Java Virtual Machine
- Automatic memory management
 - At the expense of Garbage Collection

Crash course on Java Heap

- Broken out in to sections or generations
 - Eden, or New/Young generation
 - Further sub-divided to survivors
 - Old, or Tenured generation
 - Perm, or Permanent generation
 - Disappears in Java 8
- Generations have a default size
 - Based on runtime type, client or server



Crash course on Garbage collection

- non-referenced objects are marked for deletion
- Depending on java version, generation and vm type **may** be run in parallel
- Garbage collection event types
 - Minor GC – happens often and only in Eden
 - Major GC – happens very infrequently and this Eden, Old and Perm
- In Eden
 - Objects bounce between Survivor 0 and Survivor 1
 - After surviving for so many minor garbage collections, promotion to Old generation occurs
- In Old
 - Objects are usually only removed in major (aka stop the world) garbage collection events
- In Perm
 - Objects **never** removed unless a full/major gc occurs
 - Should never see memory shrink!

Common simple tunable/run-time tweaks

Java arg	Action	Value/example	When to change
-Xmx	Sets max heap size	2g	If you really need to!
-Xms	Sets min heap size	2g	Any time
-XX:MaxPermSize	Increase permgen size	256m	If you really need to
-XX: +PrintGCDetails	Turns up logging to stdout of garbage collector	n/a	Debugging mainly but can be used full time with Xloggc
-XX: +PrintGCTimeStamps	Enhanced output of PrintGCDetails	n/a	^^
-Xloggc:<file>	Logs gc to alternate log file	/foo/bar	^^

Helpful tools

Tool Name	URL
VisualVM	https://visualvm.java.net/
Samurai	http://samuraism.jp/samurai/en/index.html
jmap	included in jdk
jstat	^^
gcviz	https://github.com/Netflix/gcviz
GCViewer	https://github.com/chewiebug/GCViewer
MAT	https://eclipse.org/mat/

Common exceptions

- `java.lang.OutOfMemoryError`: Java heap space
- `java.lang.OutOfMemoryError`: PermGen space
- `java.lang.OutOfMemoryError`: GC overhead limit exceeded
- `java.lang.NullPointerException`

How to read an exception

- When a failure occurs in java, unless the developer catches it you get a nasty backtrace
- Presented from latest to oldest
 - Eg most relevant information is on top

```

15/05/25 16:35:11 INFO DAGScheduler: Job 0 failed: count at <console>;23, took 3.016429 s
15/05/25 16:35:11 INFO TaskSetManager: Lost task 4.3 in stage 0.0 (TID 19) on executor cdm.home.bilsch.org: java.io.FileNotFoundException (/home/
vagrant/jdk-8u45-linux-x64.rpm (Permission denied)) [duplicate 19]
15/05/25 16:35:11 INFO TaskSchedulerImpl: Removed TaskSet 0.0, whose tasks have all completed, from pool
org.apache.spark.SparkException: Job aborted due to stage failure: Task 0 in stage 0.0 failed 4 times, most recent failure: Lost task 0.3 in stag
e 0.0 (TID 15, cdm.home.bilsch.org): java.io.FileNotFoundException: /home/vagrant/jdk-8u45-linux-x64.rpm (Permission denied)
    at java.io.FileInputStream.open(Native Method)
    at java.io.FileInputStream.<init>(FileInputStream.java:146)
    at org.apache.hadoop.fs.RawLocalFileSystem$LocalFSFileInputStream.<init>(RawLocalFileSystem.java:104)
    at org.apache.hadoop.fs.RawLocalFileSystem.open(RawLocalFileSystem.java:200)
    at org.apache.hadoop.fs.ChecksumFileSystem$ChecksumFSInputChecker.<init>(ChecksumFileSystem.java:141)
    at org.apache.hadoop.fs.ChecksumFileSystem.open(ChecksumFileSystem.java:341)
    at org.apache.hadoop.fs.FileSystem.open(FileSystem.java:766)
    at org.apache.hadoop.mapred.LineRecordReader.<init>(LineRecordReader.java:108)
    at org.apache.hadoop.mapred.TextInputFormat.getRecordReader(TextInputFormat.java:67)
    at org.apache.spark.rdd.HadoopRDD$$anon$1.<init>(HadoopRDD.scala:236)
    at org.apache.spark.rdd.HadoopRDD.compute(HadoopRDD.scala:212)
    at org.apache.spark.rdd.HadoopRDD.compute(HadoopRDD.scala:101)
    at org.apache.spark.rdd.RDD.computeOrReadCheckpoint(RDD.scala:277)
    at org.apache.spark.rdd.RDD.iterator(RDD.scala:244)
    at org.apache.spark.rdd.MapPartitionsRDD.compute(MapPartitionsRDD.scala:35)
    at org.apache.spark.rdd.RDD.computeOrReadCheckpoint(RDD.scala:277)
    at org.apache.spark.rdd.RDD.iterator(RDD.scala:244)
    at org.apache.spark.scheduler.ResultTask.runTask(ResultTask.scala:61)
    at org.apache.spark.scheduler.Task.run(Task.scala:64)
    at org.apache.spark.executor.Executor$TaskRunner.run(Executor.scala:203)
    at java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1145)
    at java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:615)
    at java.lang.Thread.run(Thread.java:745)

```

- Simple example, permission denied trying to read a file
- Simple version of what is going on
 - Within a java thread in org.apache.spark
 - Scheduled task trying to iterate over the contents of a file
 - Bunch of hadoop classes
 - Top most is java.io.FileInputStream in the open method

Just give me more heap!

- I used to get this request a lot!
- Keep in mind garbage collection
 - Bigger heaps mean longer garbage collections
 - But could also really just be the right answer!
- Remember other applications need memory too
 - Especially the kernel and caches!
- Having a bunch of JVMs on a single machine can be problematic when saturating host
 - Research garbage collector thread allocation!
 - <http://architects.dzone.com/articles/how-tune-java-garbage>

jstat


<

JRuby application (pid 95345)

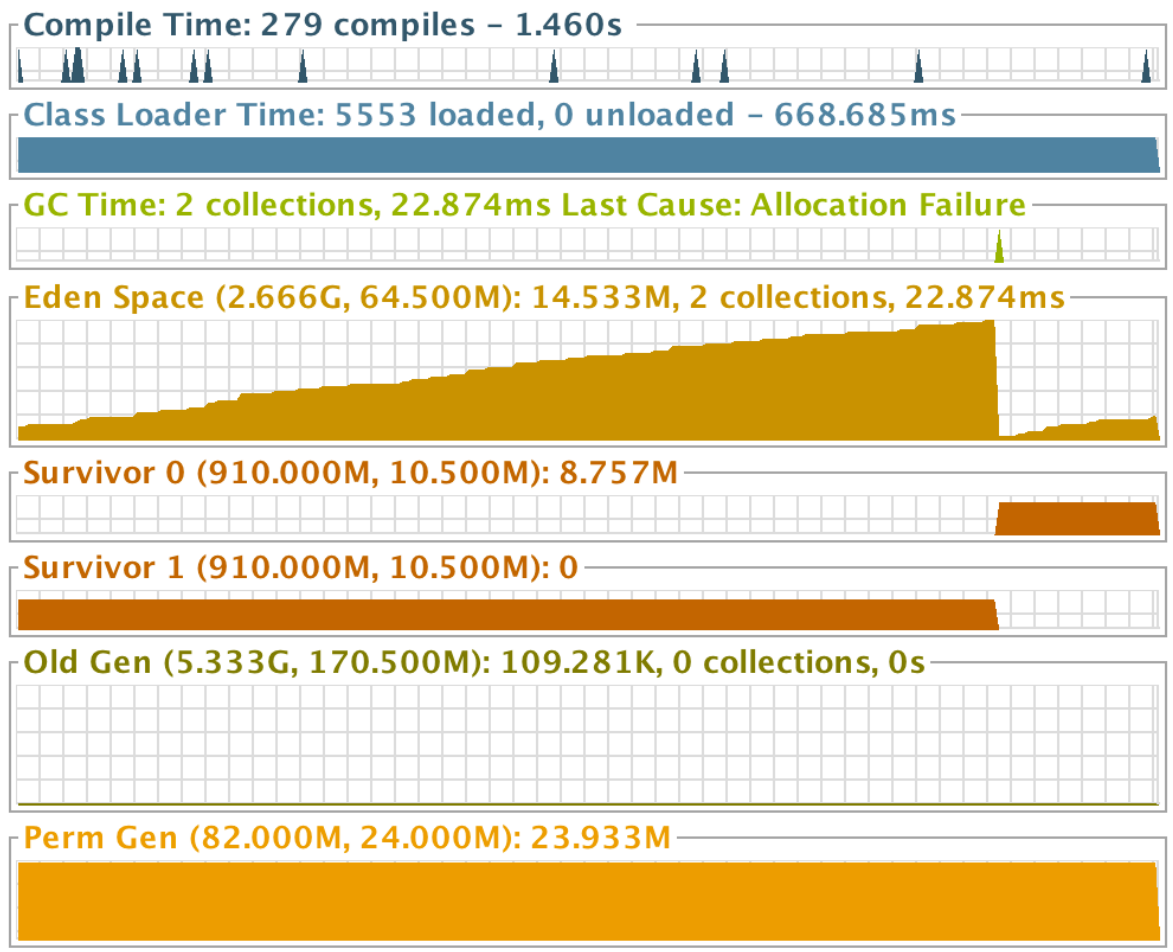
Visual GC ☒ Spaces ☒ Graphs ☐ Histogram

Refresh rate: Auto msec.

Spaces x



Graphs x




JRuby application (pid 95893)

Visual GC

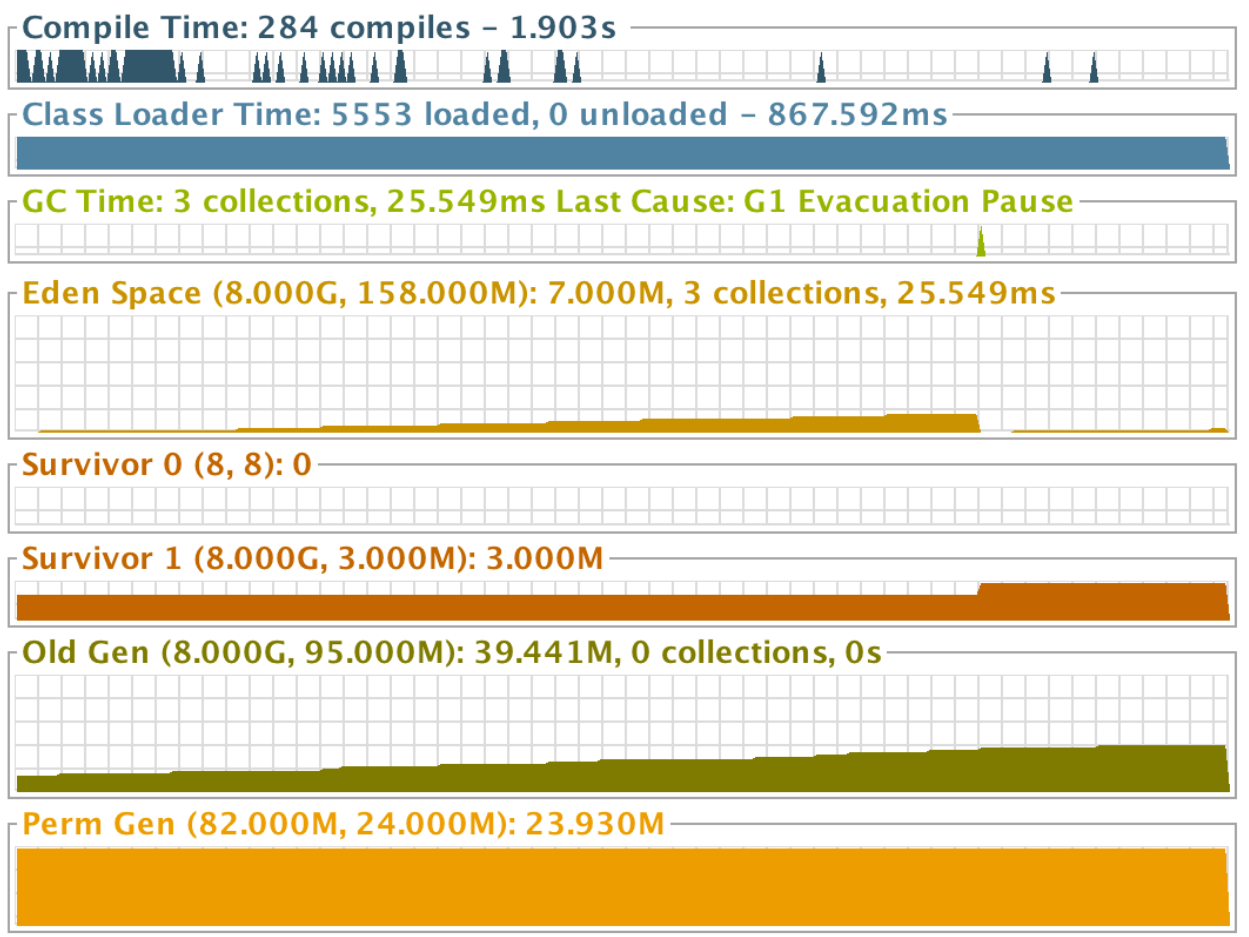
☒ Spaces ☒ Graphs ☐ Histogram

Refresh rate: Auto msec.

Spaces x



Graphs x



JMX

- Java Management Extensions
- Provides access to all kinds of neat stuff
 - Access to counters/gauges on individual MBeans
 - Ability to interrogate certain components of the MBean
 - Toggle runtime
 - Invoke methods
 - Trigger a garbage collection
 - Turn a jdbc pool on/off
 - Temporarily bump log levels
 - ...

JRuby application (pid 95893)

MBeans Browser

MBeans

- JMImplementation
- com.sun.management
- java.lang
 - ClassLoading
 - Compilation
 - GarbageCollector
 - Memory
 - MemoryManager
 - MemoryPool
 - OperatingSystem
 - Runtime
 - Threading
- java.nio
- java.util.logging

Attributes | Operations | Notifications | Metadata

Attribute values

Name	Value										
HeapMemoryUsage	<table><thead><tr><th>Name</th><th>Value</th></tr></thead><tbody><tr><td>committed</td><td>268435456</td></tr><tr><td>init</td><td>268435456</td></tr><tr><td>max</td><td>8589934592</td></tr><tr><td>used</td><td>67152288</td></tr></tbody></table>	Name	Value	committed	268435456	init	268435456	max	8589934592	used	67152288
Name	Value										
committed	268435456										
init	268435456										
max	8589934592										
used	67152288										
NonHeapMemoryUsage	javax.management.openmbean.CompositeDataSupport										
ObjectName	java.lang:type=Memory										
ObjectPendingFinalizationCount	0										
Verbose	false										

Start Page

JRuby application (pid 95345)

JRuby application (pid 95893)



Threads



Sampler



Profiler



MBeans



JConsole Plugins



Buffer Pools

JRuby application (pid 95893)

MBeans Browser

MBeans

- ▶ JMImplementation
- ▶ com.sun.management
- ▼ java.lang
 - ClassLoading
 - Compilation
 - ▶ GarbageCollector
 - Memory**
 - ▶ MemoryManager
 - ▶ MemoryPool
 - OperatingSystem
 - Runtime
 - Threading
- ▶ java.nio
- ▶ java.util.logging

Attributes

Operations

Notifications

Metadata

Operation invocation

void

gc

()

JRuby application (pid 95893)

MBeans Browser

MBeans

- ▶ JMImplementation
- ▶ com.sun.management
- ▼ java.lang
 - ClassLoading
 - Compilation
 - ▶ GarbageCollector
 - Memory
 - ▶ MemoryManager
 - ▶ MemoryPool
 - OperatingSystem
 - Runtime
 - Threading
- ▶ java.nio
- ▶ java.util.logging

Attributes | Operations | Notifications | Metadata

MBeanInfo

Name	Value
MBeanInfo	
Info:	
ObjectName	java.lang:type=Memory
ClassName	sun.management.MemoryImpl
Description	Information on the management interface of the MBean
Info Descriptor:	
immutableInfo	true
interfaceClassName	java.lang.management.MemoryMXBean
mxbean	true
MBeanAttributeInfo	
Attribute:	
Name	Verbose
Description	Verbose
Readable	true
Writable	true
Is	true
Type	boolean
Attribute Descriptor:	
openType	javax.management.openmbean.SimpleType(name=java.lang.Boolean)
originalType	boolean
MBeanAttributeInfo	
Attribute:	
Name	ObjectPendingFinalizationCount
Description	ObjectPendingFinalizationCount
Readable	true

Final thoughts, monitoring

- Last thought on JMX. Collectd and others can pull jmx metric data
 - <https://collectd.org/wiki/index.php/Plugin:GenericJMX>
 - Embeds a jvm in to collectd
 - I found upstream builds do **not** compile jmx support in
 - https://exchange.nagios.org/directory/Plugins/Java-Applications-and-Servers/check_jmx/details
 - <https://jolokia.org/index.html>
 - Requires altering the class path and may or may not be liked by vendors
 - Much lighter weight than the check_jmx as you can get at things with curl vs. invoking a small jvm